

# Dataset Description

## Features (X)

- **flength**
- **fwidth**
- **fsize**
- **fConc**
- **fConc1**
- **fAsym**
- **fM3Long**
- **fM3Trans**
- **fAlpha**
- **fDist**

## Target (y)

- **class** target class label

Imports

```
In [34]: import pandas as pd  
import numpy as np
```

`pandas` : library used for dealing with data in an csv file loading it and putting it in a pandas dataframe (2D table) having rows and columns

`numpy` : used for fast calculations mainly on arrays

```
In [35]: from sklearn.neighbors import KNeighborsClassifier  
from sklearn.preprocessing import MinMaxScaler  
import matplotlib.pyplot as plt  
from sklearn.metrics import (  
    accuracy_score,  
    precision_score,  
    recall_score,  
    confusion_matrix,  
    f1_score,  
)
```

`sklearn` : importing the KNeigboursClassifier model/ preprocessing min-max scaling (to avoid features from dominating)/metric (given target and predictions produce some values to tell how bad the model is)

# Data preprocessing

```
In [36]: # -----1. Data Preprocessing-----
# reading the data from the csv into the data frame in pandas
df = pd.read_csv("telescope_data.csv")
```

```
In [37]: df.head(1)
```

```
Out[37]:   Unnamed: 0   fLength   fWidth   fSize   fConc   fConc1   fAsym   fM3Long   fM3Trans   fAlpha
0           0    28.7967  16.0021  2.6449  0.3918  0.1982  27.7004     22.011    -8.2027   40.0
```



```
In [38]: df["class"]
```

```
Out[38]: 0      g
1      g
2      g
3      g
4      g
..
19015    h
19016    h
19017    h
19018    h
19019    h
Name: class, Length: 19020, dtype: object
```

```
In [39]: df["class"].value_counts()
```

```
Out[39]: class
g    12332
h     6688
Name: count, dtype: int64
```

syntax: df['column name'] this would return the index(by default pandas just numbered them auto increment) and will return the data of the column you specified  
df['column\_name'].value\_counts() counts the occurrences of each unique value of the column name for example here class has two values g and h so it counted all the rows that had the class=g and all the rows that has the class=h

```
In [40]: # 1. dropping the unnamed column
df.drop(columns=["Unnamed: 0"], inplace=True)
df.head(1)
```

```
Out[40]:   fLength   fWidth   fSize   fConc   fConc1   fAsym   fM3Long   fM3Trans   fAlpha   fDist
0    28.7967  16.0021  2.6449  0.3918  0.1982  27.7004     22.011    -8.2027   40.092  81.8828
```



`1.dropping the unnamed column` removed the index that was in the excel file because this is not a feature that we want to calculate the distance about we want to calculate the euclidean distance between features of points `Why do we have to remove it ?`

1. because we are calculating the distance based on the actual features of the input.
2. let's visualize having sample 1 and 1000 if distance would be  $(1000-1)^2 + \dots$  they may actually be close and similar but because the index dominates a very huge distance (it will be classified as far and the index isn't even a part of the features of the input data that we are trying to compare)

`syntax` drop helps us remove columns or rows `df.drop()` by default will remove rows so we specify we want to remove columns with the name `unnamed: 0` inplace means modify the df in place (instead of returning a new copy)

`Checking for duplicates after dropping the index column that would've made them all unique`

In [41]: `df.duplicated().sum()`

Out[41]: 115

In [42]: `df.shape`

Out[42]: (19020, 11)

In [43]: `df = df.drop_duplicates()`

In [44]: `df.shape`

Out[44]: (18905, 11)

In [45]: `df.groupby("class").head(3)`

|       | fLength  | fWidth   | fSize  | fConc  | fConc1 | fAsym    | fM3Long  | fM3Trans | fAlpha  |
|-------|----------|----------|--------|--------|--------|----------|----------|----------|---------|
| 0     | 28.7967  | 16.0021  | 2.6449 | 0.3918 | 0.1982 | 27.7004  | 22.0110  | -8.2027  | 40.0920 |
| 1     | 31.6036  | 11.7235  | 2.5185 | 0.5303 | 0.3773 | 26.2722  | 23.8238  | -9.9574  | 6.3609  |
| 2     | 162.0520 | 136.0310 | 4.0612 | 0.0374 | 0.0187 | 116.7410 | -64.8580 | -45.2160 | 76.9600 |
| 12332 | 93.7035  | 37.9432  | 3.1454 | 0.1680 | 0.1011 | 53.2566  | 89.0566  | 11.8175  | 14.1224 |
| 12333 | 102.0005 | 22.0017  | 3.3161 | 0.1064 | 0.0724 | -54.0862 | 43.0553  | -15.0647 | 88.4636 |
| 12334 | 100.2775 | 21.8784  | 3.1100 | 0.3120 | 0.1446 | -48.1834 | 57.6547  | -9.6341  | 20.7848 |

`< >`

`2.2 why do we have to use groupby` because we want to select randomly the `min_number` from both classes if we didn't use group by we would've selected randomly the

min\_value from the table in total not from both classes `groupby` just dividing the data into separate groups based on the unique values of the column name based

```
In [46]: df["class"].value_counts()
```

```
Out[46]: class
g    12332
h    6573
Name: count, dtype: int64
```

```
In [47]: df["class"].value_counts().min()
```

```
Out[47]: 6573
```

2.1 getting min of 2 classes `min class` is the h with count = 6688 will try to randomly select 6688 from both classes so that we prevent if k got greater from being biased to the majority class

```
In [48]: df.head(10)
```

```
Out[48]:   fLength  fWidth  fSize  fConc  fConc1  fAsym  fM3Long  fM3Trans  fAlpha
0    28.7967  16.0021  2.6449  0.3918  0.1982  27.7004  22.0110  -8.2027  40.0920  81.
1    31.6036  11.7235  2.5185  0.5303  0.3773  26.2722  23.8238  -9.9574  6.3609  205.
2   162.0520  136.0310  4.0612  0.0374  0.0187  116.7410  -64.8580  -45.2160  76.9600  256.
3    23.8172   9.5728  2.3385  0.6147  0.3922  27.2107  -6.4633  -7.1513  10.4490  116.
4    75.1362  30.9205  3.1611  0.3168  0.1832  -5.5277  28.5525  21.8393  4.6480  356.
5    51.6240  21.1502  2.9085  0.2420  0.1340  50.8761  43.1887  9.8145  3.6130  238.
6    48.2468  17.3565  3.0332  0.2529  0.1515  8.5730  38.0957  10.5868  4.7920  219.
7    26.7897  13.7595  2.5521  0.4236  0.2174  29.6339  20.4560  -2.9292  0.8120  237.
8   96.2327  46.5165  4.1540  0.0779  0.0390  110.3550  85.0486  43.1844  4.8540  248.
9    46.7619  15.1993  2.5786  0.3377  0.1913  24.7548  43.8771  -6.6812  7.8750  102.
```

```
In [49]: min_value = df["class"].value_counts().min()
df_balanced = df.groupby("class").sample(n=min_value, random_state=50)
df_shuffled = df_balanced.sample(frac=1, random_state=50)
df_shuffled.head(10)
```

Out[49]:

|              | fLength  | fWidth  | fSize  | fConc  | fConc1 | fAsym    | fM3Long  | fM3Trans | fAlpha  |
|--------------|----------|---------|--------|--------|--------|----------|----------|----------|---------|
| <b>8318</b>  | 54.5985  | 20.5114 | 2.7427 | 0.4286 | 0.2251 | -25.3734 | -28.8119 | -15.5386 | 3.6790  |
| <b>10955</b> | 82.0557  | 31.9178 | 3.4785 | 0.2785 | 0.1739 | 27.1880  | 39.5205  | 17.7326  | 2.2334  |
| <b>1572</b>  | 102.9340 | 36.3217 | 4.0231 | 0.1131 | 0.0584 | 51.4017  | 83.3996  | -23.4473 | 0.5982  |
| <b>16492</b> | 19.2593  | 15.2488 | 2.7701 | 0.4217 | 0.2390 | -25.5023 | 11.1513  | 13.3488  | 20.7450 |
| <b>15221</b> | 12.9176  | 11.3596 | 2.1123 | 0.7413 | 0.3900 | 15.0388  | -5.6768  | -11.5638 | 64.9330 |
| <b>552</b>   | 46.3641  | 13.7186 | 2.9763 | 0.3675 | 0.2540 | 42.3198  | 31.0888  | -10.1393 | 10.4919 |
| <b>12537</b> | 32.9878  | 8.1677  | 2.3956 | 0.5133 | 0.2637 | 38.2928  | -24.1729 | -10.2392 | 68.9474 |
| <b>16197</b> | 16.5431  | 10.4410 | 2.3801 | 0.5896 | 0.3151 | 22.2053  | 8.8734   | 12.6436  | 30.0060 |
| <b>9455</b>  | 30.3146  | 12.5905 | 2.7251 | 0.5311 | 0.3343 | 11.0797  | 27.3951  | -9.7683  | 17.1453 |
| <b>17193</b> | 24.4639  | 22.4218 | 2.6289 | 0.4371 | 0.2503 | 0.7043   | -21.2863 | 16.2589  | 87.9369 |



problem 1 here is that the indices are messed up to make it better we use .reset\_index()

In [50]:

```
min_value = df["class"].value_counts().min()
df_balanced = df.groupby("class").sample(n=min_value, random_state=50)
df_shuffled = df_balanced.sample(frac=1, random_state=50).reset_index()
df_shuffled.head(10)
```

Out[50]:

|          | index | fLength  | fWidth  | fSize  | fConc  | fConc1 | fAsym    | fM3Long  | fM3Trans | fAlpha  |
|----------|-------|----------|---------|--------|--------|--------|----------|----------|----------|---------|
| <b>0</b> | 8318  | 54.5985  | 20.5114 | 2.7427 | 0.4286 | 0.2251 | -25.3734 | -28.8119 | -15.5386 | 3.6790  |
| <b>1</b> | 10955 | 82.0557  | 31.9178 | 3.4785 | 0.2785 | 0.1739 | 27.1880  | 39.5205  | 17.7326  | 2.2334  |
| <b>2</b> | 1572  | 102.9340 | 36.3217 | 4.0231 | 0.1131 | 0.0584 | 51.4017  | 83.3996  | -23.4473 | 0.5982  |
| <b>3</b> | 16492 | 19.2593  | 15.2488 | 2.7701 | 0.4217 | 0.2390 | -25.5023 | 11.1513  | 13.3488  | 20.7450 |
| <b>4</b> | 15221 | 12.9176  | 11.3596 | 2.1123 | 0.7413 | 0.3900 | 15.0388  | -5.6768  | -11.5638 | 64.9330 |
| <b>5</b> | 552   | 46.3641  | 13.7186 | 2.9763 | 0.3675 | 0.2540 | 42.3198  | 31.0888  | -10.1393 | 10.4919 |
| <b>6</b> | 12537 | 32.9878  | 8.1677  | 2.3956 | 0.5133 | 0.2637 | 38.2928  | -24.1729 | -10.2392 | 68.9474 |
| <b>7</b> | 16197 | 16.5431  | 10.4410 | 2.3801 | 0.5896 | 0.3151 | 22.2053  | 8.8734   | 12.6436  | 30.0060 |
| <b>8</b> | 9455  | 30.3146  | 12.5905 | 2.7251 | 0.5311 | 0.3343 | 11.0797  | 27.3951  | -9.7683  | 17.1453 |
| <b>9</b> | 17193 | 24.4639  | 22.4218 | 2.6289 | 0.4371 | 0.2503 | 0.7043   | -21.2863 | 16.2589  | 87.9369 |



problem 2 now we have the index as a column that will be taken into consideration with us in the distance calculations so we have to drop it

```
In [51]: min_value = df["class"].value_counts().min()
# 2.2 groupby
# 2.3 balancing: making both classes have the same #of samples
df_balanced = df.groupby("class").sample(n=min_value, random_state=50)
# 2.4 shuffling
df_shuffled = df_balanced.sample(frac=1, random_state=50).reset_index(drop=True)
df_shuffled.head(10)
```

Out[51]:

|   | fLength  | fWidth  | fSize  | fConc  | fConc1 | fAsym    | fM3Long  | fM3Trans | fAlpha  | fl    |
|---|----------|---------|--------|--------|--------|----------|----------|----------|---------|-------|
| 0 | 54.5985  | 20.5114 | 2.7427 | 0.4286 | 0.2251 | -25.3734 | -28.8119 | -15.5386 | 3.6790  | 316.4 |
| 1 | 82.0557  | 31.9178 | 3.4785 | 0.2785 | 0.1739 | 27.1880  | 39.5205  | 17.7326  | 2.2334  | 367.0 |
| 2 | 102.9340 | 36.3217 | 4.0231 | 0.1131 | 0.0584 | 51.4017  | 83.3996  | -23.4473 | 0.5982  | 321.1 |
| 3 | 19.2593  | 15.2488 | 2.7701 | 0.4217 | 0.2390 | -25.5023 | 11.1513  | 13.3488  | 20.7450 | 88.9  |
| 4 | 12.9176  | 11.3596 | 2.1123 | 0.7413 | 0.3900 | 15.0388  | -5.6768  | -11.5638 | 64.9330 | 227.1 |
| 5 | 46.3641  | 13.7186 | 2.9763 | 0.3675 | 0.2540 | 42.3198  | 31.0888  | -10.1393 | 10.4919 | 209.6 |
| 6 | 32.9878  | 8.1677  | 2.3956 | 0.5133 | 0.2637 | 38.2928  | -24.1729 | -10.2392 | 68.9474 | 120.1 |
| 7 | 16.5431  | 10.4410 | 2.3801 | 0.5896 | 0.3151 | 22.2053  | 8.8734   | 12.6436  | 30.0060 | 214.6 |
| 8 | 30.3146  | 12.5905 | 2.7251 | 0.5311 | 0.3343 | 11.0797  | 27.3951  | -9.7683  | 17.1453 | 166.0 |
| 9 | 24.4639  | 22.4218 | 2.6289 | 0.4371 | 0.2503 | 0.7043   | -21.2863 | 16.2589  | 87.9369 | 70.7  |

◀ ➡

.sample() either takes n (a specific # of rows) or frac (frac is % of the rows) df\_balanced = df.groupby("class").sample(n=min\_value, random\_state=50) sampling after groupby -> picking randomly from both classes min\_value # of rows and the random\_state is the seed for the # random generator to produce the same results everytime it randomizes -> why to be fair if we are testing for the hyperparameter k then the accuracy is based on k not because of the change in data (and for accuracy to be the same everytime we run) df\_shuffled = df\_balanced.sample(frac=1, random\_state=50) sampling-> randomly selecting some rows from the table (100% of the table so that is mixing them up -> why is it important because if we keep them g at the top h later when splitting then most of the training data will be g while the validation and test will be h which the model never knew of and was never trained for)

random\_state why need random state if we are always taking the whole table because we need to have the same order because when we split we want to have the same sets to be fair when testing for hyperparameters

```
In [52]: df_shuffled["class"].value_counts()
```

```
Out[52]: class  
g    6573  
h    6573  
Name: count, dtype: int64
```

check if they both have the same #of samples now

Separating labels (Y) From features (X)

```
In [53]: # 3. Separate X(inputs) and Y(labels)  
X = df_shuffled.drop(columns=["class"], inplace=False)  
Y = df_shuffled["class"]
```

```
In [54]: X.head(1)
```

```
Out[54]:   fLength  fWidth  fSize  fConc  fConc1  fAsym  fM3Long  fM3Trans  fAlpha  fDis  
0    54.5985  20.5114  2.7427  0.4286  0.2251 -25.3734 -28.8119 -15.5386   3.679  316.42
```



```
In [55]: Y.head(10)
```

```
Out[55]: 0    g  
1    g  
2    g  
3    h  
4    h  
5    g  
6    h  
7    h  
8    g  
9    h  
Name: class, dtype: object
```

```
In [56]: # 4.splitting  
n = len(df_shuffled)  
# pandas expect integers it wont take floats (will not take 1.7 row for example that)  
train_end = int(0.7 * n)  
validation_end = int(0.85 * n)
```

```
In [64]: # use array slicing that is overridden by pandas  
# to be as if it is selecting rows arr[:2] then it will have rows 0 and 1  
  
# --scaling--  
scaler = MinMaxScaler()  
X_train = X[:train_end]  
X_train = scaler.fit_transform(X_train)  
Y_train = Y[:train_end]  
  
X_valid = X[train_end:validation_end]  
X_valid = scaler.transform(X_valid)  
Y_valid = Y[train_end:validation_end]
```

```
X_test = X[validation_end:]
X_test = scaler.transform(X_test)

Y_test = Y[validation_end:]
```

```
In [65]: Y_train.count()
```

```
Out[65]: 9202
```

```
In [66]: Y_test.count()
```

```
Out[66]: 1972
```

```
In [67]: Y_valid.count()
```

```
Out[67]: 1972
```

```
In [68]: # -----2.Manual knn-----
# distance
def Euc_distance(x1, x2):
    return np.sqrt(np.sum((x1 - x2) ** 2))
```

Calculating the Euclidean distance using numpy first it subtracts each corresponding elements in the 2 arrays x1,x2 then gives are result (x1[0]-x2[0]), (x1[1],x2[1]),..... and then it gets the square of each element produced as a result of subtraction -> giving an output it then will take this output and square each element in it (the output it will sum all of this together then the output -> will take a sqrt)

```
In [69]: # X could be the validation set or the test
def knn(X_train, Y_train, X, k):
    # chnaging them to numpy arrays for faster computations
    Xtrain = np.array(X_train)
    Ytrain = np.array(Y_train)
    x = np.array(X)
    Y_predicted = []
    for validation_sample in x:
        distances = []
        # defining an empty array to store all the distance between a
        # specific point and the training points
        for training_sample in Xtrain:
            distances.append(Euc_distance(validation_sample, training_sample))

        distance = np.array(distances)
        # k-nearest (returns the indicies of the nearest k elements)
        k_nearest = np.argsort(distance)[:k]

        # getting their labels
        predicted = Ytrain[k_nearest]
        # getting the labels of the k nearest (numpy feature is to pass array of in

        # majority
        labels, counts = np.unique(predicted, return_counts=True)
        majority_prediction = labels[np.argmax(counts)]
```

```

    # if the are equal the np.argmax will give us the first max
    Y_predicted.append(majority_prediction)
return Y_predicted

```

```

In [71]: # testing for multiple k values
# k < root n (which is root 13376 = 115)
kValues = [1, 3, 5, 7, 11, 13, 15, 20, 50, 60, 70, 100]
yValid = np.array(Y_valid)
accuracies = []
print("Manual knn")
for k in kValues:
    y_valid_prediction = knn(X_train, Y_train, X_valid, k)
    accuracy = np.mean(
        yValid == np.array(y_valid_prediction)
    ) # this will return an array of true and false true =1 and false =0 (it will
    print("k= ", k, "accuracy: ", accuracy)
    accuracies.append(accuracy)

best_k = kValues[np.argmax(accuracies)]
print("best k is: ", best_k)

```

```

Manual knn
k= 1 accuracy:  0.7824543610547667
k= 3 accuracy:  0.8108519269776876
k= 5 accuracy:  0.8174442190669371
k= 7 accuracy:  0.8179513184584178
k= 11 accuracy:  0.8220081135902637
k= 13 accuracy:  0.8169371196754563
k= 15 accuracy:  0.8144016227180527
k= 20 accuracy:  0.8164300202839757
k= 50 accuracy:  0.8133874239350912
k= 60 accuracy:  0.808316430020284
k= 70 accuracy:  0.8042596348884381
k= 100 accuracy:  0.800709939148073
best k is:  11

```

`scikit learn` implementation of knn algorithm , all models in knn are python object it takes a property whose name is `n_neighbors` that define the `k` (`k`-nearest neighbour)

```

In [77]: # -----3.scikit learn-----
accuracies_sk = []
for k in kValues:
    model = KNeighborsClassifier(
        n_neighbors=k
    ) # knn implementation made by scikit Learn
    y_pred = model.fit(X_train, Y_train).predict(X_valid)
    accuracy = np.mean(y_pred == Y_valid)
    accuracies_sk.append(accuracy)
    print("sklearn k=", k, "accuracy:", accuracy)

```

```
sklearn k= 1 accuracy: 0.7824543610547667
sklearn k= 3 accuracy: 0.8108519269776876
sklearn k= 5 accuracy: 0.8174442190669371
sklearn k= 7 accuracy: 0.8179513184584178
sklearn k= 11 accuracy: 0.8220081135902637
sklearn k= 13 accuracy: 0.8169371196754563
sklearn k= 15 accuracy: 0.8144016227180527
sklearn k= 20 accuracy: 0.8164300202839757
sklearn k= 50 accuracy: 0.8133874239350912
sklearn k= 60 accuracy: 0.808316430020284
sklearn k= 70 accuracy: 0.8042596348884381
sklearn k= 100 accuracy: 0.800709939148073
```

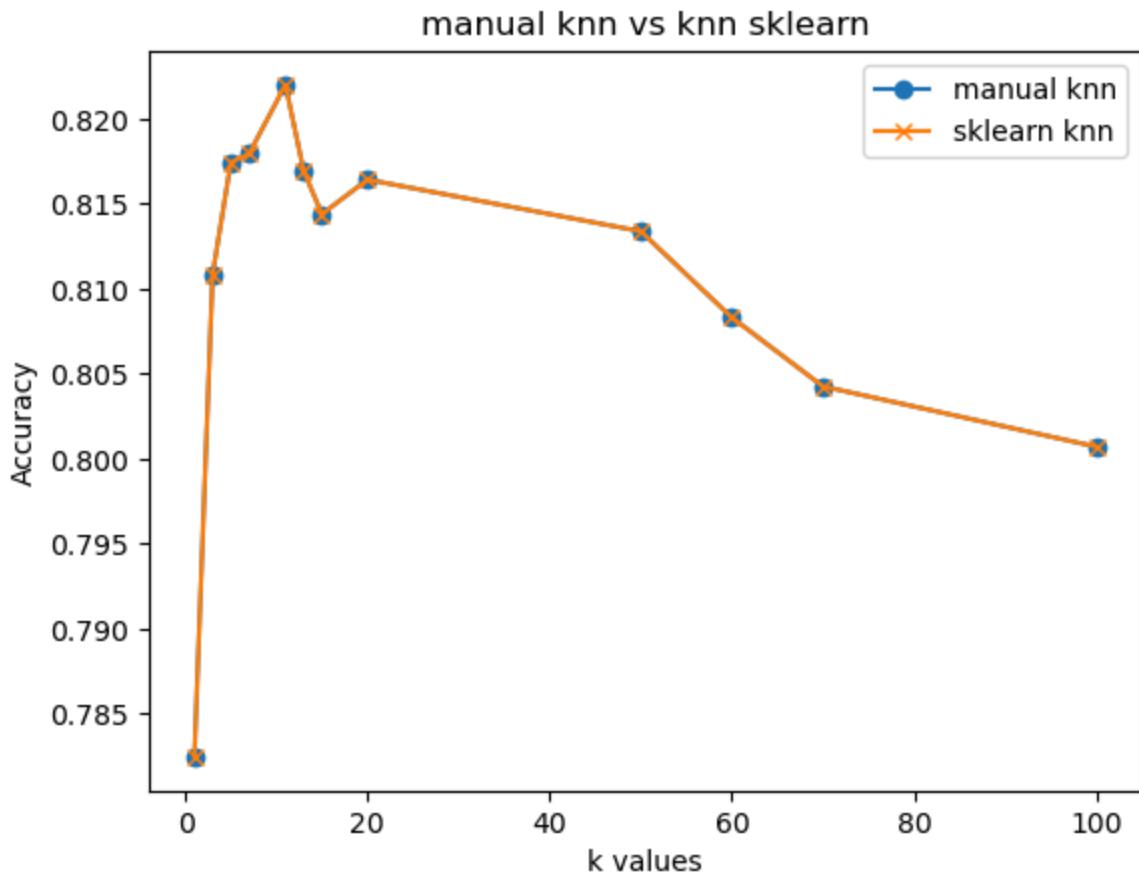
.fit(X\_train,Y\_train) -> the model learns from the training data **Note** KNN is a 'lazy learner' so all it does is just stores the training points it doesn't do any training computations

.predict() predict the values of the X\_validation (After getting trained on the test data)

**What it does?**

- compute all distances to all the training points
- find the k nearest neighbors
- pick the class by majority vote giving us the y\_pred which is the prediction of the validation set now we can compare with the Y\_valid which is the correct label to know the accuracy of our model

```
In [78]: # -----plot accuracy vs k values -----
plt.plot(kValues, accuracies, label="manual knn", marker="o")
plt.plot(kValues, accuracies_sk, label="sklearn knn", marker="x")
plt.xlabel("k values")
plt.ylabel("Accuracy")
plt.title("manual knn vs knn sklearn")
plt.legend() # to show the Labels (manual knn and skLearn)
plt.show()
```



## Predicting the test data after tuning

starting off with manual knn

```
In [79]: # -----prediction the test set after tuning-----
y_test_prediction = np.array(knn(X_train, Y_train, X_test, best_k))
ytest = np.array(Y_test)
```

calculating metric manually

```
In [80]: # ----- Accuracy of manual -----
accuracy_manual = np.mean(Y_test == y_test_prediction)
```

```
In [81]: # ----- getting the confusion matrix (TP,FP,FN,TN)-----
# predicted positive and it is actually positive
# assuming g is the positive class
TP = np.sum(
    (y_test_prediction == "g") & (ytest == "g")
) # bitwise and to take each element in the array (not try and evaluate the whole
# predicted to be negative class = h and it is actually negative
TN = np.sum((y_test_prediction == "h") & (ytest == "h"))
# predicted positive and it is actually negative
FP = np.sum((y_test_prediction == "g") & (ytest == "h"))
FN = np.sum((y_test_prediction == "h") & (ytest == "g"))
```

```
In [99]: print("manual knn:")
print("accuracy ", accuracy_manual)
print("confusion matrix: ")
print(f"[{TP} {FN}]")
print(f"[{FP} {TN}]")
precision = TP / (TP + FP) # out of all positives what was actually positive
print("precision of manual knn: ", precision)
recall = TP / (TP + FN) # out of all actual positives hm got predicted correctly
print("recall for manual knn: ", recall)
fscore = (2 * (precision * recall)) / (precision + recall)
print("fscore of manual knn: ", fscore)
```

```
manual knn:
accuracy 0.815922920892495
confusion matrix:
[849 124]
[239 760]
precision of manual knn: 0.7803308823529411
recall for manual knn: 0.8725590955806783
fscore of manual knn: 0.8238719068413392
```

```
In [93]: # -----metrics using scikit learn built metric
y_test_prediction_sklearn =
    KNeighborsClassifier(n_neighbors=11).fit(X_train, Y_train).predict(X_test)
```

```
In [100...]: print("For sklearn: ")
print("accuracy: ", accuracy_score(ytest, y_test_prediction_sklearn))
print("confusion matrix: ", confusion_matrix(ytest, y_test_prediction_sklearn))
print("precision: ", precision_score(ytest, y_test_prediction_sklearn, pos_label="g"))
print("recall: ", recall_score(ytest, y_test_prediction_sklearn, pos_label="g"))
print("fscore: ", f1_score(ytest, y_test_prediction_sklearn, pos_label="g"))

For sklearn:
accuracy: 0.815922920892495
confusion matrix: [[849 124]
 [239 760]]
precision: 0.7803308823529411
recall: 0.8725590955806783
fscore: 0.8238719068413392
```

## 6. Comments on the results

# RESULTS

```
Manual knn
k= 1 accuracy: 0.7824543610547667
k= 3 accuracy: 0.8108519269776876
k= 5 accuracy: 0.8174442190669371
k= 7 accuracy: 0.8179513184584178
k= 11 accuracy: 0.8220081135902637
k= 13 accuracy: 0.8169371196754563
k= 15 accuracy: 0.8144016227180527
```

```

k= 20 accuracy: 0.8164300202839757
k= 50 accuracy: 0.8133874239350912
k= 60 accuracy: 0.808316430020284
k= 70 accuracy: 0.8042596348884381
k= 100 accuracy: 0.800709939148073
best k is: 11
sklearn k= 1 accuracy: 0.7824543610547667
sklearn k= 3 accuracy: 0.8108519269776876
sklearn k= 5 accuracy: 0.8174442190669371
sklearn k= 7 accuracy: 0.8179513184584178
sklearn k= 11 accuracy: 0.8220081135902637
sklearn k= 13 accuracy: 0.8169371196754563
sklearn k= 15 accuracy: 0.8144016227180527
sklearn k= 20 accuracy: 0.8164300202839757
sklearn k= 50 accuracy: 0.8133874239350912
sklearn k= 60 accuracy: 0.808316430020284
sklearn k= 70 accuracy: 0.8042596348884381
sklearn k= 100 accuracy: 0.800709939148073

```

- at k = 1 the model overfits which means that it is too specific to the data set so it can't generalize well (it is sensitive to noise)
- it starts getting better till the k=11
- it starts decreasing again because now the model starts to underfit (model is too simple most probably it may classify a point to the other class) it may be a point very far from the point we are trying to predict but because we are taking a huge k it will still be classified as that far point (majority)

Comparison observation: Manual and sklearn implementations produce identical results (same accuracy for each k value), confirming the correctness of the manual implementation. Why are they identical because they are using the same approach for calculating the distance and majority

| Metric    | KNN Manual | Sklearn KNN |
|-----------|------------|-------------|
| Accuracy  | 0.815      | 0.815       |
| Precision | 0.78       | 0.78        |
| Recall    | 0.87       | 0.87        |
| F1-Score  | 0.82       | 0.82        |